

SIMPLE LINEAR REGRESSION

Our task is to predict the student score based on the number hours that he had spent on study. We had used simple linear regression which is the supervised regression technique, to predict the score of the student

```
In [1]: 1 #step 1: import all required libraries
        2 import pandas as pd
        3 import numpy as np
        4 import matplotlib.pyplot as plt
        5 from sklearn.model_selection import train_test_split
        6 from sklearn.linear_model import LinearRegression
        7 from sklearn.metrics import mean_absolute_error, r2_score
```

```
c:\users\vamsi2001\appdata\local\programs\python\python39\lib\site-packages\num
py\_distributor_init.py:30: UserWarning: loaded more than 1 DLL from .libs:
c:\users\vamsi2001\appdata\local\programs\python\python39\lib\site-packages\num
py\.libs\libopenblas.EL2C6PLE4ZYW3ECEVIV30XXGRN2NRFM2.gfortran-win_amd64.dll
c:\users\vamsi2001\appdata\local\programs\python\python39\lib\site-packages\num
py\.libs\libopenblas.XWYDX2IKJW2NMTWSFYNGFUWKQU3LYTCZ.gfortran-win_amd64.dll
warnings.warn("loaded more than 1 DLL from .libs:")
```

```
In [2]: 1 #step 2: Load the data
        2 data=pd.read_csv('marks.csv')
        3 data.head()
```

Out[2]:

	Hours	Scores
0	3.75	67.79
1	9.51	99.06
2	7.32	86.43
3	5.99	74.11
4	1.56	63.51

```
In [3]: 1 data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype
---  -
0   Hours   149 non-null         float64
1   Scores  148 non-null         float64
dtypes: float64(2)
memory usage: 2.5 KB
```

```
In [5]: 1 data.isnull().sum()
```

```
Out[5]: Hours      1  
Scores      2  
dtype: int64
```

```
In [6]: 1 # Handling Missing values  
2 data['Hours']=data['Hours'].fillna(data['Hours'].mean())  
3 data['Scores']=data['Scores'].fillna(data['Scores'].median())
```

```
In [7]: 1 data.isnull().sum()
```

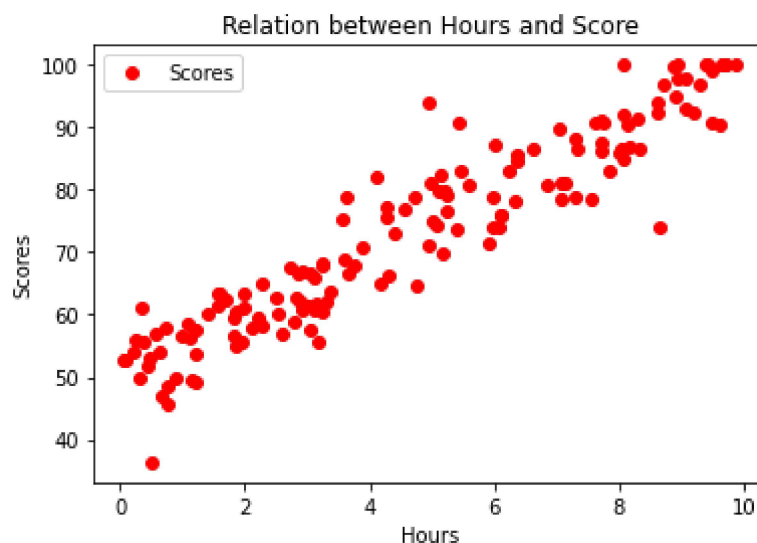
```
Out[7]: Hours      0  
Scores      0  
dtype: int64
```

```
In [8]: 1 data.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 150 entries, 0 to 149  
Data columns (total 2 columns):  
#   Column  Non-Null Count  Dtype  
---  -  
0   Hours    150 non-null    float64  
1   Scores   150 non-null    float64  
dtypes: float64(2)  
memory usage: 2.5 KB
```

```
In [9]: 1 #step 3:plot the relationship between features(columns)  
2 data.plot(x='Hours',y='Scores',style='o',color='red')  
3 plt.title('Relation between Hours and Score')  
4 plt.xlabel('Hours')  
5 plt.ylabel('Scores')
```

```
Out[9]: Text(0, 0.5, 'Scores')
```



```
In [10]: 1 #step 4: split dataset
2 X=data.iloc[:, :-1].values
3 y=data.iloc[:, 1].values
4 #X=np.array(data['Hours']).reshape(-1,1)
5 #y=np.array(data['Scores'])
6 X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.2)
```

```
In [11]: 1 #step 5: build the model
2 reg=LinearRegression()
3 reg.fit(X_train,y_train)
```

Out[11]: LinearRegression()

```
In [12]: 1 #step 6:Predictions
2 y_pred=reg.predict(X_test)
3 y_pred
```

Out[12]: array([77.05107029, 86.08942867, 84.86802889, 63.95766464, 56.18956204,
 94.05295524, 76.85564633, 55.99413807, 85.16116484, 65.8630483 ,
 87.26197246, 61.51486508, 53.50248252, 65.03249645, 74.90140668,
 75.19454263, 80.76412563, 59.36520147, 73.38687095, 63.56681671,
 97.52173062, 70.35779949, 89.50934805, 92.39185154, 52.13451477,
 66.84016813, 90.97502779, 85.99171669, 64.25080059, 75.29225461])

```
In [13]: 1 print("m value",reg.coef_)#m value
2 print("Intercept value",reg.intercept_)# intercept
```


m value [4.88559912]
Intercept value 50.3268430917942

```
In [14]: 1 #step 7:Model Evaluation
2 print("mean absolute error is",mean_absolute_error(y_test,y_pred))
3 print("r2 value",r2_score(y_test,y_pred))
```

mean absolute error is 4.67946019193863
r2 value 0.8561505991497251

```
In [15]: 1 #step 8:plot Regression Line
2 line=reg.coef_*X+reg.intercept_ #y=mx+c
3 plt.scatter(X,y)
4 plt.plot(X,line,c='red')
```

Out[15]: [<matplotlib.lines.Line2D at 0x1e562da75b0>]



```
In [16]: 1 #step 9:predict using custom input
2 hr=np.array(8)# hours
3 pred=reg.predict(hr.reshape(-1,1))
4 print(pred)
```

[89.41163607]

```
In [ ]: 1
```

```
In [ ]: 1
```