

Our task is to Heartfailure prediction using SVM classifier

Here is the description about the dataset

age: The age of the patient.

anaemia: Presence of anemia in the patient (1 = Yes, 0 = No).

creatinine_phosphokinase: Level of the enzyme creatinine phosphokinase in the blood (measured in mcg/L).

diabetes: Whether the patient has diabetes (1 = Yes, 0 = No).

ejection_fraction: Percentage of blood leaving the heart with each contraction (measured in percentage).

high_blood_pressure: Whether the patient has high blood pressure (1 = Yes, 0 = No).

platelets: Number of platelets in the blood (measured in kiloplatelets/mL).

serum_creatinine: Level of serum creatinine in the blood (measured in mg/dL).

serum_sodium: Level of serum sodium in the blood (measured in mEq/L).

sex: The biological sex of the patient (1 = Male, 0 = Female).

smoking: Whether the patient smokes (1 = Yes, 0 = No).

time: Follow-up period for the patient (measured in days).

DEATH_EVENT: The target variable indicating if the patient died during the follow-up period (1 = Yes, 0 = No).

In [2]:

```
1 #Step1:importing
2 import pandas as pd
3 import numpy as np
4 import matplotlib.pyplot as plt
5 import seaborn as sns
6 from sklearn.preprocessing import StandardScaler
7 from sklearn.model_selection import train_test_split
8 from sklearn.metrics import accuracy_score, confusion_matrix, classification_r
```

```
c:\users\vamsi2001\appdata\local\programs\python\python39\lib\site-packages\nu
py\_distributor_init.py:30: UserWarning: loaded more than 1 DLL from .libs:
c:\users\vamsi2001\appdata\local\programs\python\python39\lib\site-packages\nu
py\.libs\libopenblas.EL2C6PLE4ZYW3ECEVIV3OXXGRN2NRFM2.gfortran-win_amd64.dll
c:\users\vamsi2001\appdata\local\programs\python\python39\lib\site-packages\nu
py\.libs\libopenblas.XWYDX2IKJW2NMTWSFYNGFUWKQU3LYTCZ.gfortran-win_amd64.dll
warnings.warn("loaded more than 1 DLL from .libs:")
```

```
In [3]: 1 df=pd.read_csv('HeartFailure.csv')
        2 df.head()
```

```
Out[3]:
```

	age	anaemia	creatinine_phosphokinase	diabetes	ejection_fraction	high_blood_pressure	platelets
0	75.0	0	582	0	20	1	26500
1	55.0	0	7861	0	38	0	26335
2	65.0	0	146	0	20	0	16200
3	50.0	1	111	0	20	0	21000
4	65.0	1	160	1	20	0	32700

```
In [4]: 1 df.info()
        2 df.describe()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 299 entries, 0 to 298
Data columns (total 13 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   age                                   299 non-null    float64
1   anaemia                               299 non-null    int64
2   creatinine_phosphokinase              299 non-null    int64
3   diabetes                              299 non-null    int64
4   ejection_fraction                     299 non-null    int64
5   high_blood_pressure                    299 non-null    int64
6   platelets                             299 non-null    float64
7   serum_creatinine                      299 non-null    float64
8   serum_sodium                          299 non-null    int64
9   sex                                   299 non-null    int64
10  smoking                               299 non-null    int64
11  time                                  299 non-null    int64
12  DEATH_EVENT                           299 non-null    int64
dtypes: float64(3), int64(10)
memory usage: 30.5 KB
```

```
Out[4]:
```

	age	anaemia	creatinine_phosphokinase	diabetes	ejection_fraction	high_blood_
count	299.000000	299.000000	299.000000	299.000000	299.000000	29
mean	60.833893	0.431438	581.839465	0.418060	38.083612	
std	11.894809	0.496107	970.287881	0.494067	11.834841	
min	40.000000	0.000000	23.000000	0.000000	14.000000	
25%	51.000000	0.000000	116.500000	0.000000	30.000000	
50%	60.000000	0.000000	250.000000	0.000000	38.000000	
75%	70.000000	1.000000	582.000000	1.000000	45.000000	
max	95.000000	1.000000	7861.000000	1.000000	80.000000	

In []:

1

In [34]:

1 df.columns

```
Out[34]: Index(['age', 'anaemia', 'creatinine_phosphokinase', 'diabetes',
               'ejection_fraction', 'high_blood_pressure', 'platelets',
               'serum_creatinine', 'serum_sodium', 'sex', 'smoking', 'time',
               'DEATH_EVENT'],
              dtype='object')
```

In [5]:

```
1 cat_cols=['anaemia','diabetes','high_blood_pressure','sex','smoking','DEATH_
2 for i in cat_cols:
3     df[i]=df[i].astype('category')
4 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 299 entries, 0 to 298
```

```
Data columns (total 13 columns):
```

#	Column	Non-Null Count	Dtype
0	age	299 non-null	float64
1	anaemia	299 non-null	category
2	creatinine_phosphokinase	299 non-null	int64
3	diabetes	299 non-null	category
4	ejection_fraction	299 non-null	int64
5	high_blood_pressure	299 non-null	category
6	platelets	299 non-null	float64
7	serum_creatinine	299 non-null	float64
8	serum_sodium	299 non-null	int64
9	sex	299 non-null	category
10	smoking	299 non-null	category
11	time	299 non-null	int64
12	DEATH_EVENT	299 non-null	category

```
dtypes: category(6), float64(3), int64(4)
```

```
memory usage: 19.0 KB
```

In [6]:


1 df.describe()

Out[6]:

	age	creatinine_phosphokinase	ejection_fraction	platelets	serum_creatinine	se
count	299.000000	299.000000	299.000000	299.000000	299.000000	
mean	60.833893	581.839465	38.083612	263358.029264	1.39388	
std	11.894809	970.287881	11.834841	97804.236869	1.03451	
min	40.000000	23.000000	14.000000	25100.000000	0.50000	
25%	51.000000	116.500000	30.000000	212500.000000	0.90000	
50%	60.000000	250.000000	38.000000	262000.000000	1.10000	
75%	70.000000	582.000000	45.000000	303500.000000	1.40000	
max	95.000000	7861.000000	80.000000	850000.000000	9.40000	


```
In [7]: 1 plt.figure(figsize=(10,5))
        2 corrmat=df.corr()
        3 sns.heatmap(corrmat, annot=True)
```

Out[7]: <AxesSubplot:>



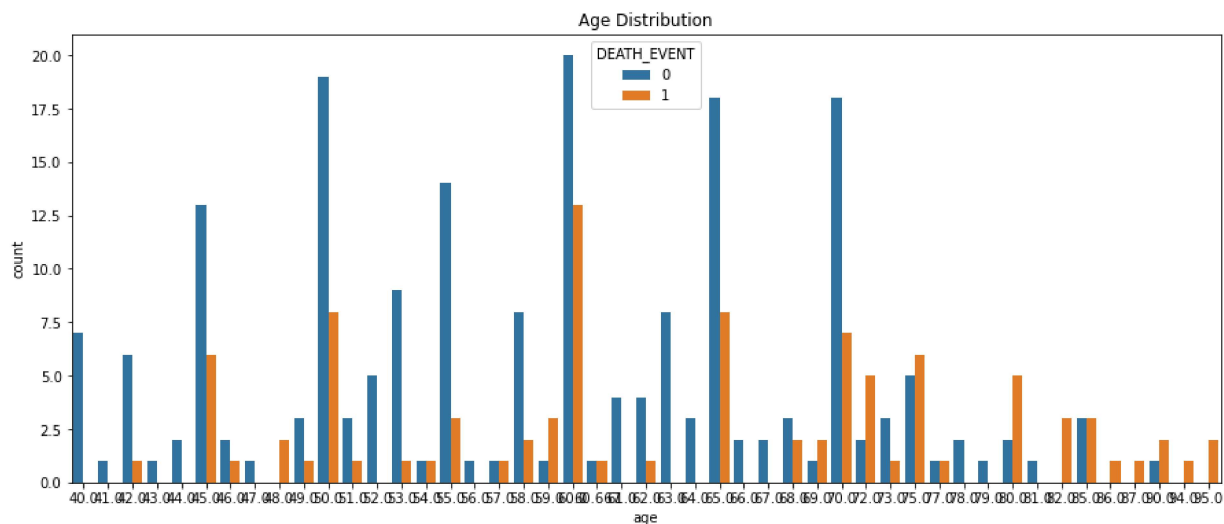
```
In [8]: 1 event_counts=df['DEATH_EVENT'].value_counts()
2
3 plt.figure(figsize=(8,5))
4 ax=sns.countplot(x=df['DEATH_EVENT'])
5 ax.bar_label(ax.containers[0])
```

Out[8]: [Text(0, 0, '203'), Text(0, 0, '96')]



```
In [9]: 1 plt.figure(figsize=(15,6))
2 ax=sns.countplot(x=df['age'], data=df, hue=df['DEATH_EVENT'])
3 ax.set_title('Age Distribution')
```

Out[9]: Text(0.5, 1.0, 'Age Distribution')



```
In [10]: 1 X=df.drop(columns='DEATH_EVENT', axis=1)
2 Y=df['DEATH_EVENT']
```

```
In [11]: 1 cols=list(X.columns)
2 scaler=StandardScaler()
3 X_scaled=scaler.fit_transform(X)
4 X_scaled_df=pd.DataFrame(X_scaled, columns=cols)
5 X_scaled_df.describe().T
```

```
Out[11]:
```

	count	mean	std	min	25%	50%	75%
age	299.0	5.703353e-16	1.001676	-1.754448	-0.828124	-0.070223	0.771889
anaemia	299.0	1.009969e-16	1.001676	-0.871105	-0.871105	-0.871105	1.147968
creatinine_phosphokinase	299.0	0.000000e+00	1.001676	-0.576918	-0.480393	-0.342574	0.000166
diabetes	299.0	9.060014e-17	1.001676	-0.847579	-0.847579	-0.847579	1.179830
ejection_fraction	299.0	-3.267546e-17	1.001676	-2.038387	-0.684180	-0.007077	0.585389
high_blood_pressure	299.0	0.000000e+00	1.001676	-0.735688	-0.735688	-0.735688	1.359272
platelets	299.0	7.723291e-17	1.001676	-2.440155	-0.520870	-0.013908	0.411120
serum_creatinine	299.0	1.425838e-16	1.001676	-0.865509	-0.478205	-0.284552	0.005926
serum_sodium	299.0	-8.673849e-16	1.001676	-5.363206	-0.595996	0.085034	0.766064
sex	299.0	-8.911489e-18	1.001676	-1.359272	-1.359272	0.735688	0.735688
smoking	299.0	-1.188199e-17	1.001676	-0.687682	-0.687682	-0.687682	1.454161
time	299.0	-1.901118e-16	1.001676	-1.629502	-0.739000	-0.196954	0.938759

```
In [42]: 1 X_scaled_df.head()
```

```
Out[42]:
```

	age	anaemia	creatinine_phosphokinase	diabetes	ejection_fraction	high_blood_pressure
0	1.192945	-0.871105	0.000166	-0.847579	-1.530560	1.359272
1	-0.491279	-0.871105	7.514640	-0.847579	-0.007077	-0.735688
2	0.350833	-0.871105	-0.449939	-0.847579	-1.530560	-0.735688
3	-0.912335	1.147968	-0.486071	-0.847579	-1.530560	-0.735688
4	0.350833	1.147968	-0.435486	1.179830	-1.530560	-0.735688

```
In [12]: 1
2 x_train, x_test, y_train, y_test = train_test_split(X_scaled_df, Y, test_size=0.2)
```

```
In [14]: 1 from sklearn.svm import SVC
2 modelsvm=SVC(kernel='sigmoid').fit(x_train, y_train)
```

```
In [15]: 1 y_pred=modelsvm.predict(x_test)
          2 y_pred
```


```
Out[15]: array([0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
                0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0,
                1, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 1, 1,
                0, 1, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1,
                1, 0]), dtype=int64)
```

```
In [16]: 1 print(accuracy_score(y_test, y_pred))
          2
```

```
0.7444444444444445
```

```
In [17]: 1 sns.heatmap(confusion_matrix(y_test, y_pred),annot=True,cmap='Blues')
          2
```

```
Out[17]: <AxesSubplot:>
```



```
In [18]: 1 print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.72	0.92	0.81	53
1	0.82	0.49	0.61	37
accuracy			0.74	90
macro avg	0.77	0.71	0.71	90
weighted avg	0.76	0.74	0.73	90

GridSearchCV is a hyperparameter tuning technique provided by scikit learn.

It exhaustively searches through all possible combinations of hyperparameter values you specify.

For each combination, it evaluates model performance using crossvalidation.

Finally, it picks the combination that gives the best performance.

```
In [49]: 1 from sklearn.model_selection import GridSearchCV
2
3 param_grid = {
4     'C': [0.1, 1, 10, 100],
5     'gamma': ['scale', 'auto', 0.01, 0.1, 1],
6     'kernel': ['rbf', 'linear', 'poly']
7 }
8
9 grid = GridSearchCV(SVC(), param_grid, cv=5, scoring='accuracy')
10 grid.fit(x_train, y_train)
11
12 print("Best Parameters:", grid.best_params_)
13
14 best_model = grid.best_estimator_
15 y_pred = best_model.predict(x_test)
16
17 print("Accuracy after tuning:", accuracy_score(y_test, y_pred))
18
```

Best Parameters: {'C': 10, 'gamma': 0.01, 'kernel': 'rbf'}
Accuracy after tuning: 0.7777777777777778

C: Regularization parameter which controls the trade-off between correct classification and margin size.

gamma: Kernel coefficient (for rbf, poly, sigmoid). It controls how far the influence of a single training example reaches.

kernel: The type of kernel function to use (linear, rbf, or poly).

This creates $4 \times 5 \times 3 = 60$ total combinations to try

In []:

1

In []:

1

In []:

1

In []:

1